

Deep Learning for NLP

When Big Data hits Machine Learning



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Nils Reimers

Course-Website: www.deeplearning4nlp.com

Organization of this Course

- This course will **not** give a long motivation on deep learning
- This course will focus on *hands-on exercise*, less on theory
 - There are tons of good introductions and lectures that cover all the math
 - We approach this topic more like an engineer:
 - Use the things that work
 - Deep knowledge only needed when something doesn't work
- Weekly meeting, Monday at 11am
 - Please read the **preparation for class** papers— it will be assumed that you know the content
 - Theoretical input, clarification of questions for the papers etc.
 - Practice tasks (programming)
 - For a given task (NER, Genre Classification, Sentiment Classification) we develop a deep neural network using Python / Theano / Lasagne

Recommended Readings

- Recommended Readings for this lecture:
 - From the [2009 Yoshua Bengio Book](#):
 - *1 Introduction* - Good introduction into the terminology and the basic concept of deep learning
 - *2 Theoretical Advantages of Deep Architectures* - In case you are interested why deep architectures are better than shallow
 - [Chapter 1 - Using neural nets to recognize handwritten digits](#)
 - [Chapter 2 - How the backpropagation algorithm works](#)
 - From the [2015 Yoshua Bengio Book](#):
 - *1 Introduction* - if you are interested in the historical development of neural networks
 - *Part I: Applied Math and Machine Learning Basics* - As a refresher for linear algebra and machine learning basics (if needed)
 - *6 Feedforward Deep Networks* - Read on feedforward networks, the most simple type of neural networks

What is Deep Learning?



- Deep Learning is a subfield of Machine Learning
- Most machine learning methods work well because of human-designed features
 - Recasens et al. (2009) lists >30 features for event coreference resolution
 - Most time spent on engineering features
- Machine learning becomes optimizing weights to make a final prediction given the features

Feature	Definition
PRON_m1	m1 is a pronoun
PRON_m2	m2 is a pronoun
HEAD_MATCH	Head match
WORDNET_MATCH	EuroWordNet match
NP_m1	m1 NP type
NP_m2	m2 NP type
NE_m1	m1 NE type
NE_m2	m2 NE type
NE_MATCH	NE match
SUPERTYPE_MATCH	Supertype match
GENDER_AGR	Gender agreement
NUMBER_AGR	Number agreement
ACRONYM	m2 is an acronym of m1
QUOTES	m2 is in quotes
FUNCTION_m1	m1 function
FUNCTION_m2	m2 function
COUNT_m1	m1 count
COUNT_m2	m2 count
SENT_DIST	Sentence distance
MENTION_DIST	Mention distance
WORD_DIST	Mention distance

Machine Learning vs. Deep Learning

Traditional Machine Learning

Feature Engineering:
Describe your data with features a computer can
understand

Machine Learning:
Some hyper-
parameter tuning

- Task, language and text-domain specific
- Requires high domain expertise (Ph.D.-level)
- Often requires other tools and resources (POS-tagger, Wordnet ...)

Deep Learning Approach

Getting
Domain
Expertise

Design / select a suitable
network architecture

Optimize architecture & fine-tune
parameters

Deep Learning

- *Representation Learning* attempts to automatically learn good features or representations
- *Deep Learning* attempts to learn multiple levels of representation from **raw input**



- Instead of modeling the problem by the design of features, deep learning attempts to learn by itself a good representation (i.e. features).
- Large amounts of data are typically needed
 - They must not necessarily be labeled (see autoencoders, word embeddings)
 - For small datasets: Hand-designed features can still be included

The Deep Learning Dream

- Input the pure characters and output the desired label (sentiment, genre, translated version, QA)
- No hassle with preprocessing:
 - No error propagation
 - No inconsistencies between pipeline components (different segmentations, tagsets etc.)
 - No 'component not available in your language'
- This makes models easily adaptable to new text domains and languages
 - Just take existent network architecture and train on new dataset
- Dream not yet achieved, but we are on a good way
 - See: Zang et al., 2015, *Text Understanding from Scratch*
 - Zang et al., 2015, *Character-level Convolutional Networks for Text Classification will be published on NIPS 2015*

What is New?

- Deep Neural Networks are nothing new (started ~1960)
- Until 2006 training of deep neural networks wasn't successful
 - Hinton et al. (2006) presented a new, stacked training method which applies the concept of pre-training
- Successful training of deep networks (from scratch) requires large amounts of data
- Training requires lots of computation power
 - GPUs are perfect to train deep neural networks
 - Training is complex, but at inference time deep neural networks are often superior

Neural Network Basics



TECHNISCHE
UNIVERSITÄT
DARMSTADT




Neural Network Basics

- Given several **inputs**: $x_1, x_2, x_3, \dots \in \mathbb{R}$
and several **weights**: $w_1, w_2, w_3, \dots \in \mathbb{R}$
and a **bias** value: $b \in \mathbb{R}$

- A neuron produces a single output:


$$o_1 = s(\sum_i w_i x_i + b)$$

- This sum $\sum_i w_i x_i + b$ is called the **activation** of the neuron
- The function s is called the **activation function** for the neuron
- The weights and bias values are typically initialized randomly and learned during training



Activation Function & Non-linearity


- The critical part is the non-linearity
- Example: Create a classifier, that learns the XOR function
- No linear classifier can solve this problem



- “not bad” \neq “not” + “bad”


Activation Function

- The non-linearity is a crucial concept that gives neural networks more representational power compared to some other techniques (linear SVM, logistic regression)
- Without the non-linearity, it is impossible to model certain combinations of features (like Boolean XOR function), unless we do manual feature engineering
- Typical non-linear activation functions for neural nets are the *sigmoid function*, the *hyperbolic tangent* or a *step function*



Typical Activation Function


- Sigmoid Function scales between 0 and 1



Img-Source: Wikipedia

Typical Activation Function

- *hyperbolic tangent* scales between -1 and 1



Img-Source: Wikipedia

Typical Activation Function


- Rectifier (blue):

$$f(x) = \max(0, x)$$

- Smooth approximation (green):

$$f(x) = \ln(1 + e^x)$$

- A unit using the rectifier function is called rectified linear unit (ReLU)



Img-Source: Wikipedia

Which Activation Function to use?


- Most papers use the *tanh* or a mathematically easier to compute version of it (due to performance increases)
- For certain problems the rectifier can be useful, especially for convolutional neural networks
 - Be careful with your input data, the activation can become arbitrarily large
-> *“one input dimension could rule them all”*

Rule of Thumb

- Use *tanh* for all your hidden layers
- Changing the activation function has only minor effect



Feed Forward Neural Networks



Img-Source: <http://ufldl.stanford.edu/wiki/>

Feed Forward Networks

- The most simple network architecture
- Information flows only in one direction
- Input Layer: Represents my data
- Output layer: Represents possible labels

- The hidden layer (L_2 , L_3) represent learned non-linear combination of input data
- For solving the XOR problem, we need a hidden layer
 - some neurons in the hidden layer will activate only for some combination of input features
 - the output layer can represent combination of the activations of the hidden neurons
- Neural network with one hidden layer **is a universal approximator**
 - Every function can be modeled as a shallow feed forward network
 - Not all functions can be represented *efficiently* with a single hidden layer
⇒ we still need deep neural networks


Matrix Notation

- Given 3-dimensional input $x \in \mathbb{R}^3$
- Given first weight matrix: $W_1 \in \mathbb{R}^{3 \times 3}$
- Given first bias vector: $b_1 \in \mathbb{R}^3$

- Given second weight matrix: $W_2 \in \mathbb{R}^{3 \times 2}$
- Given second bias vector: $b_2 \in \mathbb{R}^2$

- Given third weight matrix: $W_3 \in \mathbb{R}^{2 \times 2}$
- Given third bias vector: $b_3 \in \mathbb{R}^2$

- Computation L_2 : $l_2 = \tanh(W_1 x + b_1)$
- Computation L_3 : $l_3 = \tanh(W_2 l_2 + b_2)$
- Computation L_4 : $l_4 = \text{softmax}(W_3 l_3 + b_3)$



Img-Source: <http://ufldl.stanford.edu/wiki/>

- For classification, most networks apply a softmax-classifier as final layer
- Softmax regression is a generalization of logistic regression to the case of multiple classes
- Given we have K classes (K = number of output units). Compute the activation z for the last layer:


$$z = W_3 l_3 + b_3 \in \mathbb{R}^K$$

Compute the final output y :

$$y_j = \text{softmax}(z)_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

- y_j can have values between 0 and 1
- y_j sums up to 1 -> It can be interpreted as probability distribution

Feed-Forward Network for Handwritten Digit Recognition



Img-Source: <http://neuralnetworksanddeeplearning.com>

Rule of Thumb

- The number and the size of the hidden layers can have a large impact on the performance
- More hidden layers \Rightarrow more parameters to learn \Rightarrow more data needed
- Start with a small number of hidden layers, i.e. with 1
- Increase number of hidden layers stepwise until you find an optimum
- Typically decreasing sizes for the hidden layers, for example
2000 dim \Rightarrow 1500 dim \Rightarrow 750 dim \Rightarrow 100 dim \Rightarrow 10 dim



Initialization of Weights and Bias

- The weights and bias vectors are the parameters that we learn during training
- The bias is typically initialized to 0
- The weights are randomly initialized – Use Glorot-style uniform initialization
 - Xavier Glorot & Yoshua Bengio, *Understanding the difficulty of training deep feedforward neural networks*
(<http://jmlr.org/proceedings/papers/v9/glorot10a/glorot10a.pdf>)
 - Explanation of the algorithm:
<http://andyljones.tumblr.com/post/110998971763/an-explanation-of-xavier-initialization>
 - Implementation / Usage in Theano & Lasagne is straight forward
 - Be careful: The maximal size of the weights is different for the *tanh* and the *sigmoid* activation function.

Training Neural Networks



TECHNISCHE
UNIVERSITÄT
DARMSTADT



Img-Source: <https://www.flickr.com/photos/jakerust/>

Defining an Error Function

- First we need an error function which computes the difference between the output of the network and the expected output (true label)

- In case of single label classification, we can use the negative log-likelihood:

$$E(x, W, b) = -\log(o_y)$$


with y denoting the expected label and o the output vector.

- In case of a distributional classification, we can use the mean squared error:

$$E(x, W, b) = \frac{1}{2} \sum (y_i - o_i)^2$$


with y_i denoting the expected value for node i

- We then want to minimize the error function
- Defining the error function is an **important aspect** in designing a deep neural network
 - Changing the error function changes what our network learns



Img-Source: Wikipedia

Back Propagation – Gradient Descent



Source: mathworks.de

Gradient Descent

- We want to minimize the error function by tuning the trainable parameters (weights and biases)
- The gradient (the derivative of a multi dim. function) points us towards a local minima
- We follow the gradient by a certain step length. This length is called *learning rate*

Computation of Gradients in Theano

- Computation of the gradient (the derivative of a multi dimensional function) can be cumbersome
 - Especially as the computation must be computationally efficient, as we will perform billions of such computations
- Theano is great, as it provides us automatic gradient computation
 - We don't need to compute the derivative
 - And we don't need to implement it into our program code
 - -> Awesome!
- Next slides look a bit complicated, but doing back propagation via stochastic gradient descent (SGD) is quite simple in Theano

Back Propagation – Gradient Descent

Training with back propagation:


1. Given the input data, we compute the values for the output neurons
2. Compare the output to the gold labels – compute error function
3. Compute the derivative for all tunable parameters (weights and parameters)
4. Update the parameters:

$$W^{(i)} := W^{(i)} - \lambda \frac{\partial}{\partial W^{(i)}} E(x, W, b)$$

$$b^{(i)} := b^{(i)} - \lambda \frac{\partial}{\partial b^{(i)}} E(x, W, b)$$

λ is denoting the learning rate

5. Each iteration is called an *epoch*. With each epoch, the error function decreases, converging to a local minima




Source: Mikolov, 2014

Motivation for Deep Learning

- Feed forward networks with a single hidden layer can compute any function \Rightarrow in theory, no need for deep architectures
- However, learning shallow architectures is not always efficient
- To learn the parity function (N bits at input, output is 1 if the number of active input bits is odd) requires an exponential number of hidden units \Rightarrow requires exponentially more training data
- See the 2009 Yoshua Bengio Book – *Chapter 2 Theoretical Advantages of Deep Architectures for more information*
 - http://www.iro.umontreal.ca/%7Ebengioy/papers/ftml_book.pdf

Going from Shallow to Deep Neural Networks

- Neural Networks can have several hidden layers
- Initializing the weights randomly and training all layers at once does hardly work
- Instead we train layerwise on unannotated data (a.k.a. pre-training):
 - Train the first hidden layer
 - Fix the parameters for the first layer and train the second layer.
 - Fix the parameters for the first & second layer, train the third layer
 - ...
- After the pre-training, train all layers using your annotated data
- The pre-training on your unannotated data creates a high-level abstractions of the input data
- The final training with annotated data fine tunes all parameters in the network




Everything is a Vector



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Beautiful Dance Moves



Source: <https://www.flickr.com/photos/dylanng/>

Vectors in Deep Learning

- Deep Learning loves dense vectors
- Everything is represented and understood as a vector
- Vectors allows an end-to-end-training
- No more in-between mapping to tag sets etc.
 - No problem with modeling hard cases and ambiguities
 - No problem with error propagation between pipeline components

Dense Vector

- In a dense vector most values are non-zero




Representations at NLP Levels: Morphology

- Traditional: Morphemes

prefix	stem	suffix
un	interest	ed

- DL:
 - every morpheme is a vector
 - a neural network combines two vectors into one vector
 - Thang et al. 2013




Source: Richard Socher, CS224d, <http://cs224d.stanford.edu/syllabus.html>

Neural word vectors - visualization




Representations at NLP Levels: Syntax

- Traditional: Phrases
Discrete categories like NP, VP



- DL:
 - Every word and every phrase is a vector
 - a neural network combines two vectors into one vector
 - Socher et al. 2011

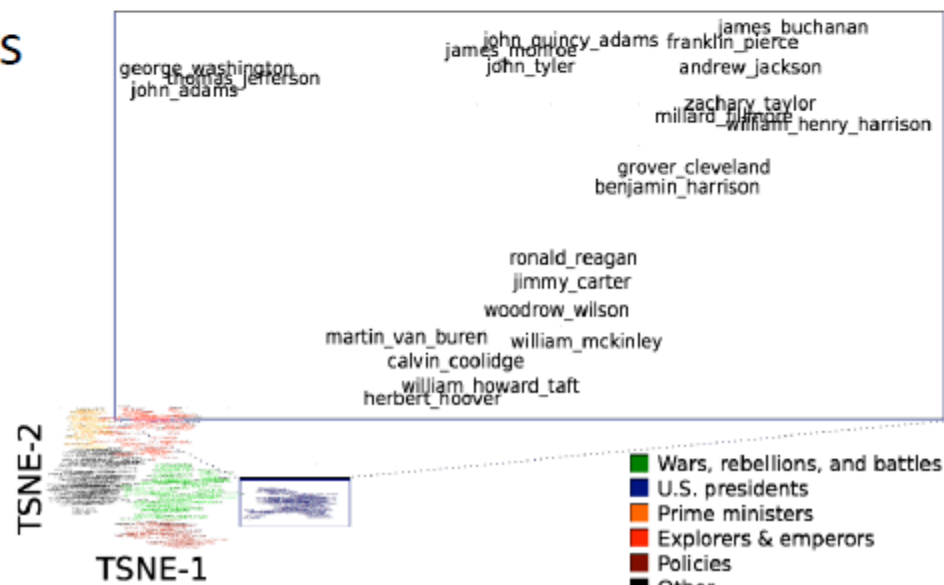


Question Answering

- Common: A lot of feature engineering to capture world and other knowledge, e.g. regular expressions, Berant et al. (2014)


Is main verb trigger?			
Yes		No	
Condition	Regular Exp.	Condition	Regular Exp.
Wh- word subjective?	AGENT	default	(ENABLE SUPER) ⁺
Wh- word object?	THEME	DIRECT	(ENABLE SUPER)
		PREVENT	(ENABLE SUPER)*PREVENT(ENABLE SUPER)*

- DL: Same deep learning model that was used for morphology, syntax, logical semantics and sentiment can be used!
- Facts are stored in vectors

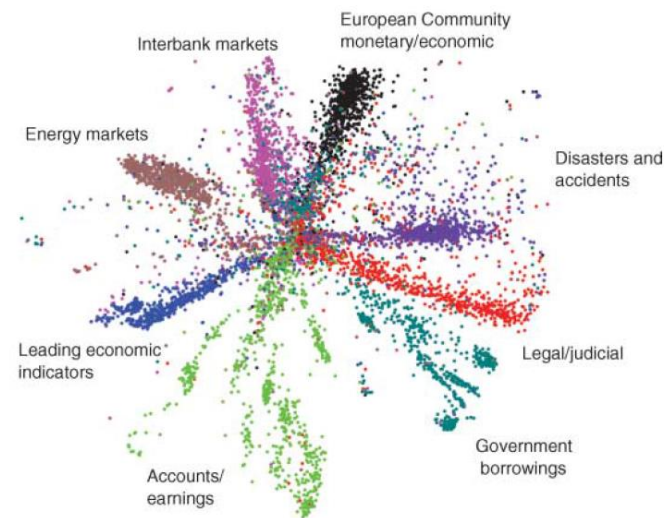


Learning Multiple Levels of Representation

- Biologically inspired: The brain has a deep architecture
- Instead of solving a task directly, we learn intermediate representations
- Highly non-linear properties can be captured this way
- Insufficient model depth can be exponentially inefficient



LSA



Deep Autoencoder

Source: Hinton et al., 2006

Next Lecture

■ Preparation before class:

- Install Python (2.7), NumPy, SciPy and Theano ([Installing Theano for Ubuntu](#))
- Install [Lasagne](#)
- Refresh your knowledge on Python and Numpy:
 - [Python and Numpy Tutorial](#)
 - [Python-Tutorial](#) and [Numpy refresher](#) from the Theano website
- **Hint:** You can install Python, Theano etc. on you local desktop machine and log into it via SSH or via [IPython Notebook](#) during class