

# Deep Learning for NLP

## Convolutional Neural Networks



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

Nils Reimers

**Course-Website: [www.deeplearning4nlp.com](http://www.deeplearning4nlp.com)**

# Recommended Readings

- <https://www.youtube.com/watch?v=EevTPpQvxiU>
- Kim, 2014, *Convolutional Neural Networks for Sentence Classification*
- <http://cs231n.github.io/convolutional-networks/>

# Convolutional Neural Network

- Universal architecture achieving state-of-the-art performance

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-rand	76.1	45.0	82.7	89.6	91.2	79.8	83.4
CNN-static	81.0	45.5	86.8	93.0	92.8	84.7	<b>89.6</b>
CNN-non-static	<b>81.5</b>	48.0	87.2	93.4	93.6	84.3	89.5
CNN-multichannel	81.1	47.4	<b>88.1</b>	93.2	92.2	<b>85.0</b>	89.4
RAE (Socher et al., 2011)	77.7	43.2	82.4	—	—	—	86.4
MV-RNN (Socher et al., 2012)	79.0	44.4	82.9	—	—	—	—
RNTN (Socher et al., 2013)	—	45.7	85.4	—	—	—	—
DCNN (Kalchbrenner et al., 2014)	—	48.5	86.8	—	93.0	—	—
Paragraph-Vec (Le and Mikolov, 2014)	—	<b>48.7</b>	87.8	—	—	—	—
CCAE (Hermann and Blunsom, 2013)	77.8	—	—	—	—	—	87.2
Sent-Parser (Dong et al., 2014)	79.5	—	—	—	—	—	86.3
NBSVM (Wang and Manning, 2012)	79.4	—	—	93.2	—	81.8	86.3
MNB (Wang and Manning, 2012)	79.0	—	—	<b>93.6</b>	—	80.0	86.3
G-Dropout (Wang and Manning, 2013)	79.0	—	—	93.4	—	82.1	86.1
F-Dropout (Wang and Manning, 2013)	79.1	—	—	<b>93.6</b>	—	81.9	86.3
Tree-CRF (Nakagawa et al., 2010)	77.3	—	—	—	—	81.4	86.1
CRF-PR (Yang and Cardie, 2014)	—	—	—	—	—	82.7	—
SVM <sub>S</sub> (Silva et al., 2011)	—	—	—	—	<b>95.0</b>	—	—

Kim, 2014. Performance on Sentence Classification Tasks

# Some notes on Conv. Neural Networks

- Convolutional Neural Networks are dominating computer vision
- For NLP, they became popular in ~2014
- Understanding the notation is quite difficult in computer vision
  - Images are typically 3 dimensional (width x height x color)
  - In NLP, we mainly deal with 1 dimensional data (our sentence / document)
  - Notation for 1 dimensional data is much simpler

# Convolutional Neural Networks solve 2 crucial Challenges

## ▪ First Challenge:

- In a lot of cases, we have variable sized input data, e.g. length of sentence / length of document.
- Our network / our hidden layers are of fixed sizes
- Solution 1: Window Approach (Collobert et al.), but we don't capture information outside of the window
- Solution 2: Recursive & Recurrent Neural Networks

## ▪ Second Challenge:

- Increasing the window in the window approach allows us to capture more context information, but increases dramatically the number of parameters
- Often the position in the context window is of minor importance:
  - Jim [sells]<sub>Pred</sub> his car for [\$5,000]<sub>???</sub>
  - Jim [sells]<sub>Pred</sub> his car, which he inherited from his dad, for [\$5,000]<sub>???</sub>

# Single Layer CNN – Single Filter

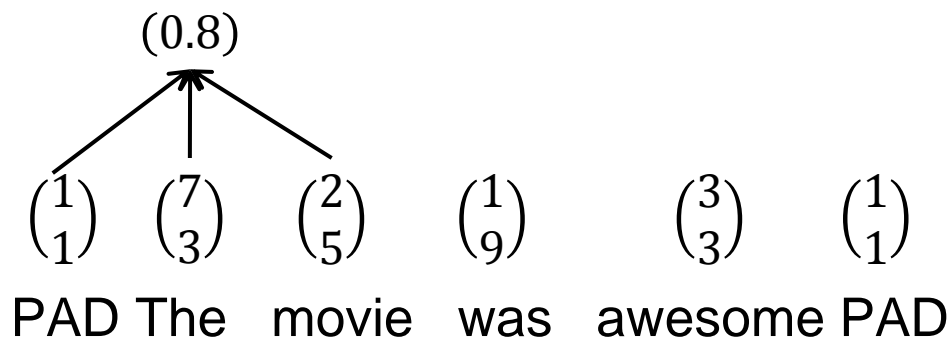
- We compute a single filter for a window size of  $n$  (here  $n=3$ ):

Word Vectors:  $w_i \in \mathbb{R}^2$

Weight Matrix:  $W \in \mathbb{R}^{1 \times 6}$

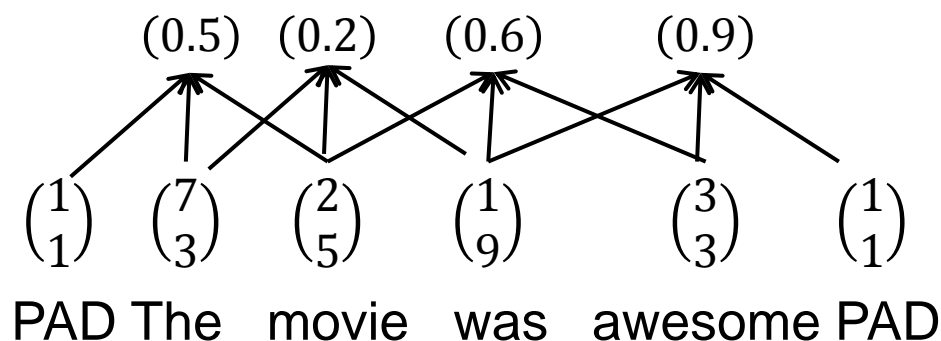
Bias:  $b \in \mathbb{R}$

$$\text{output} = \tanh \left( W \begin{pmatrix} w_1 \\ w_2 \\ w_3 \end{pmatrix} + b \right)$$



# Single Layer CNN – Single Filter

- Compute the output for all windows of size  $n$  (in our case  $n=3$ )
- For each window use the same weight and bias values (shared weights)
- This gives us the same number of digits as the length of the sentence



# Single Layer CNN – Pooling Layer

- New building block: Pooling
- Idea: Capture the most important activation
- Let  $o_1, o_2, \dots \in \mathbb{R}$  denote the output values for our filter
- Compute a max-over-time pooling layer:


$$c = \max_i(o_i) \in \mathbb{R}$$

- Because of max-over-time pooling, length of input sequence is irrelevant.
- We could use the output  $c$  and forward it to a softmax classifier and derive a sentiment class for the sentence
- Max-pooling most common in NLP. In Computer Vision, min-pooling and mean-pooling also common.



# Excursion: Max Pooling in Computer Vision

- In computer vision, pooling is often applied over fixed windows (e.g. 2x2)
- Be careful, don't confuse max pooling and max-over-time pooling



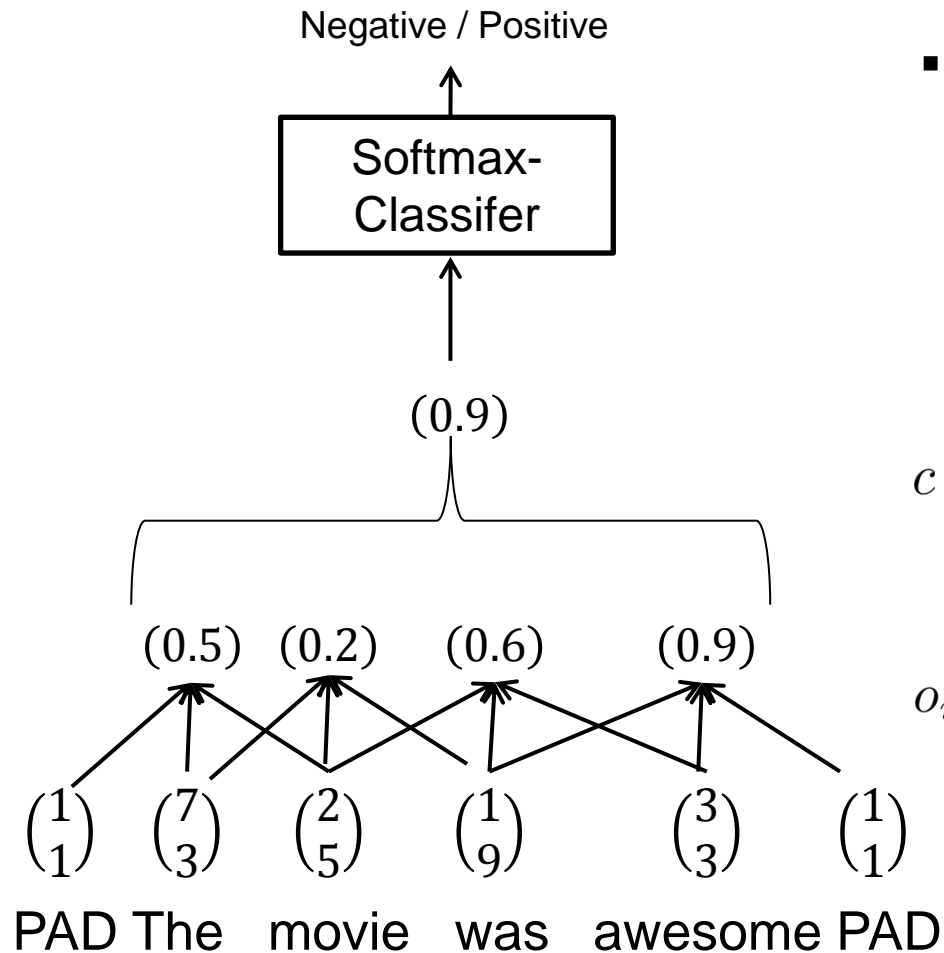
## Example:

- 224x224 pixel gray scale images, 64 images / batch
- 2x2 max-pooling reduces each image to 112 x 112 dimensions

# Max Pooling vs. Max-over-Time

- Max pool with 2x2 filters on variable sized input generates variable sized output
- Max-over-time generates fixed-sized output
- In NLP, mostly max-over-time is used  
(as presented by Collobert et al., NLP almost from scratch, section 3.2.2)
- A lot of literature on max pooling originated from computer vision
  - Be careful with different terminology and hyper parameters for those pooling layers
- Most libraries (Lasagne, Keras) are optimized for computer vision and only support window sized max pooling
  - Using the existent layers for 1-dimensional max-over-time can be a bottleneck
  - Implementation of max-over-time quite easy in Theano: *T.max(matrix, axis=1)*

# Single Layer CNN + Classification

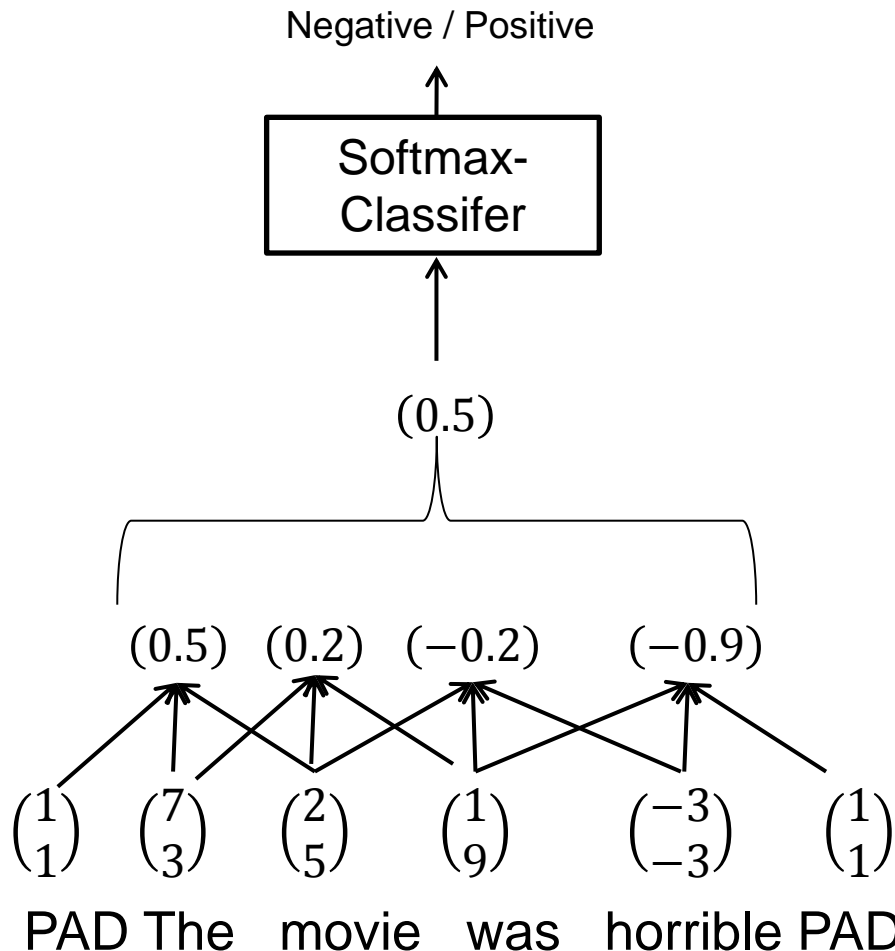


- You can train this like any other Neural Network

$$c = \max_i(o_i)$$

$$o_i = \tanh \left( W \begin{pmatrix} w_{i-1} \\ w_i \\ w_{i+1} \end{pmatrix} + b \right)$$

# 1 Dimension – not enough information

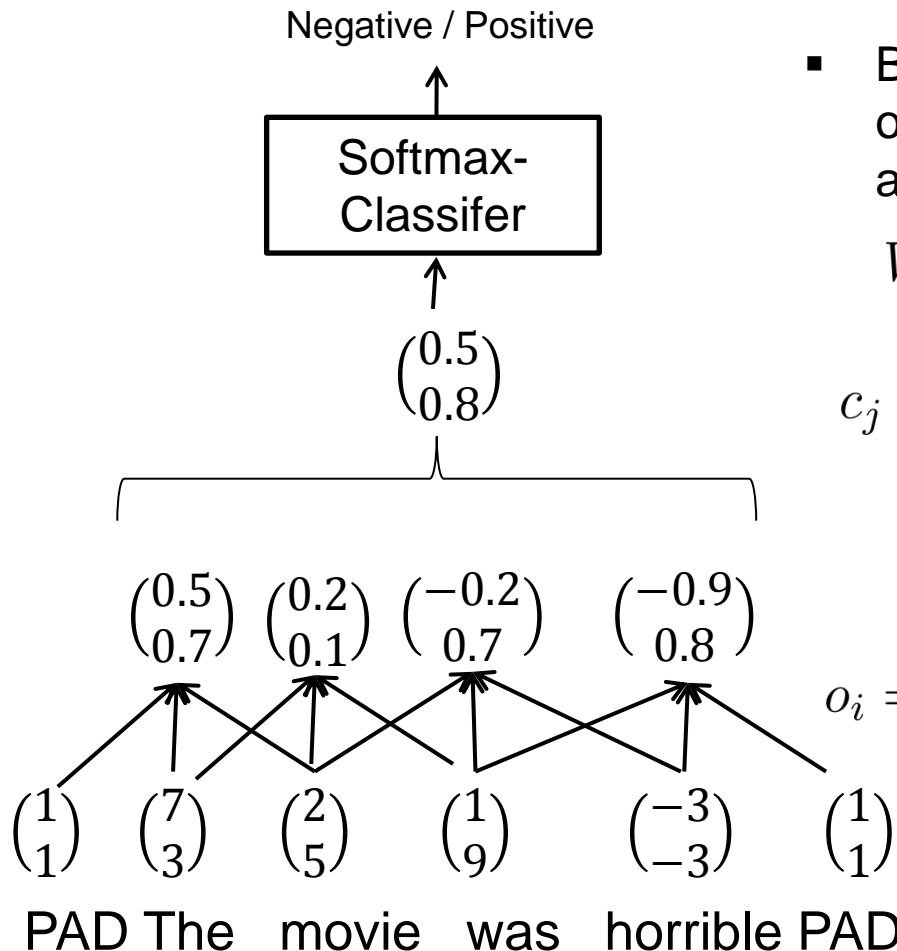


- With only a single filter our possibilities are limited

$$c = \max_i(o_i)$$

$$o_i = \tanh \left( W \begin{pmatrix} w_{i-1} \\ w_i \\ w_{i+1} \end{pmatrix} + b \right)$$

# Single Layer CNN – Multiple Filters



- By changing the dimensionality of the weight matrix, we can add further filters:

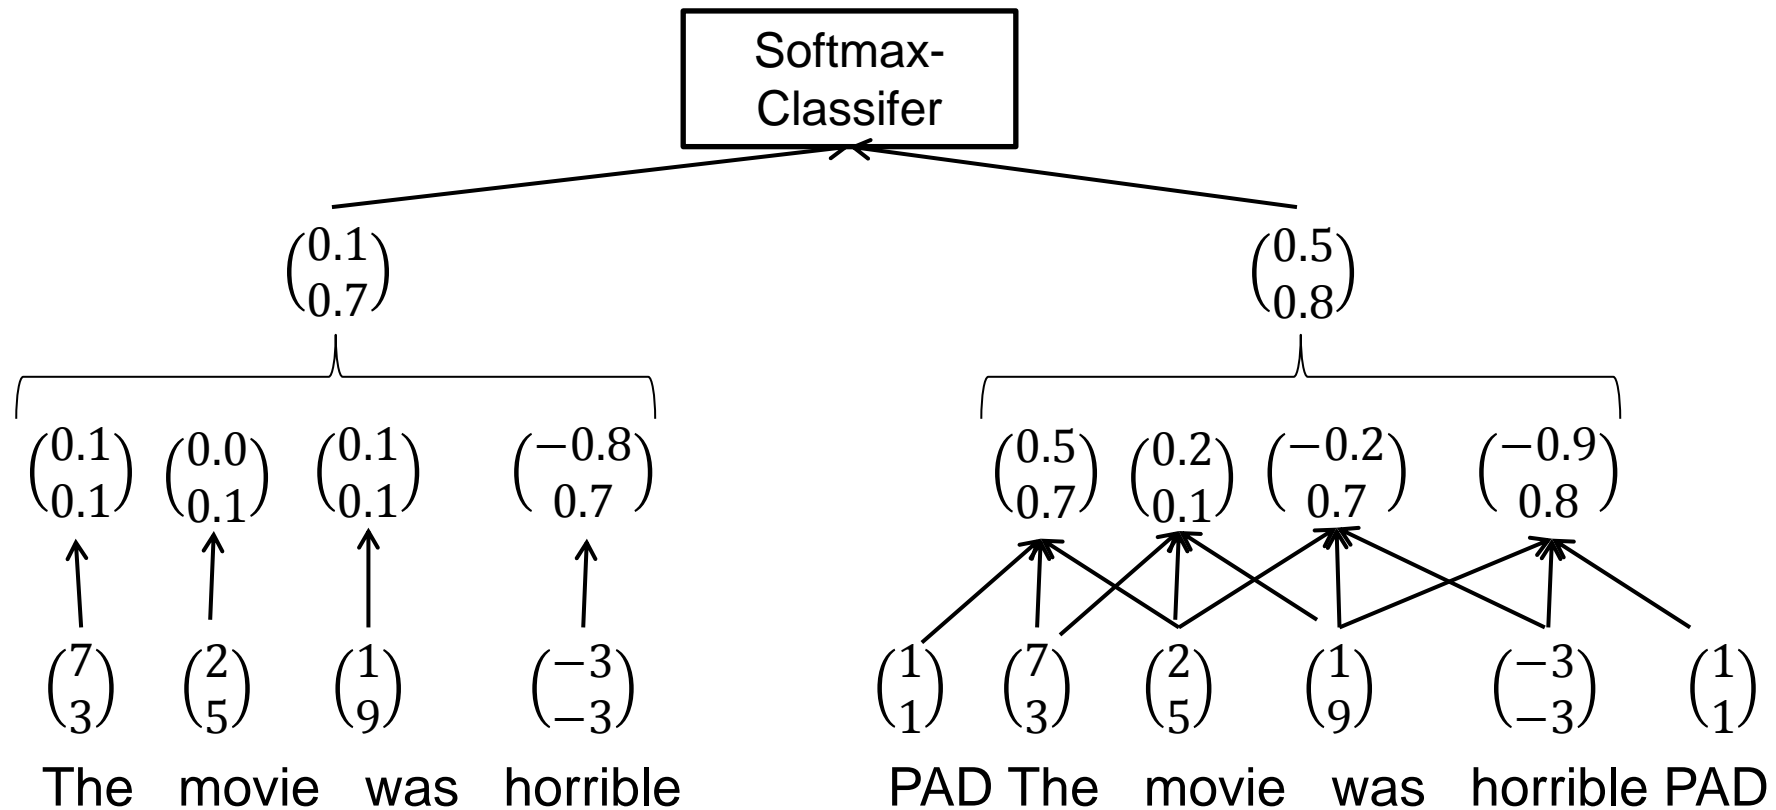
$$W \in \mathbb{R}^{k \times 6}$$

$$c_j = \max_i(o_{i,j}) \text{ for } 0 < j < k$$

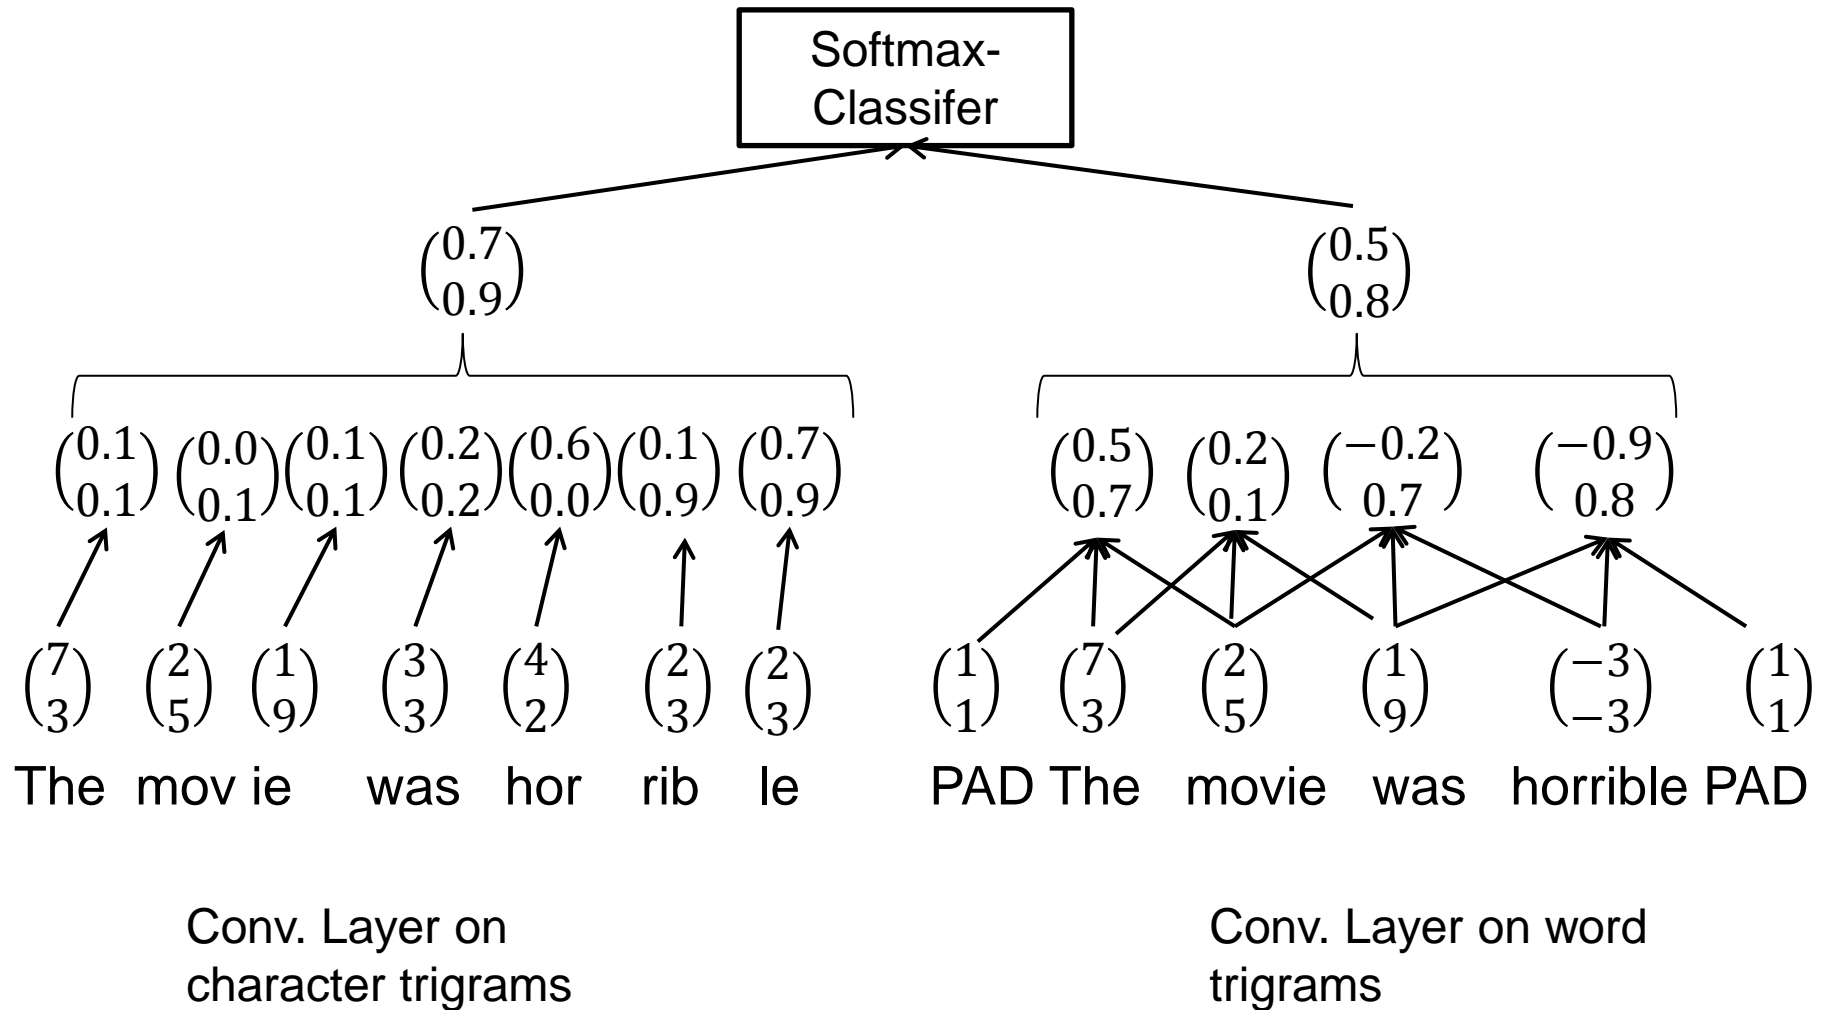
$$o_i = \tanh \left( W \begin{pmatrix} w_{i-1} \\ w_i \\ w_{i+1} \end{pmatrix} + b \right) \in \mathbb{R}^k$$

# Going further with Filters

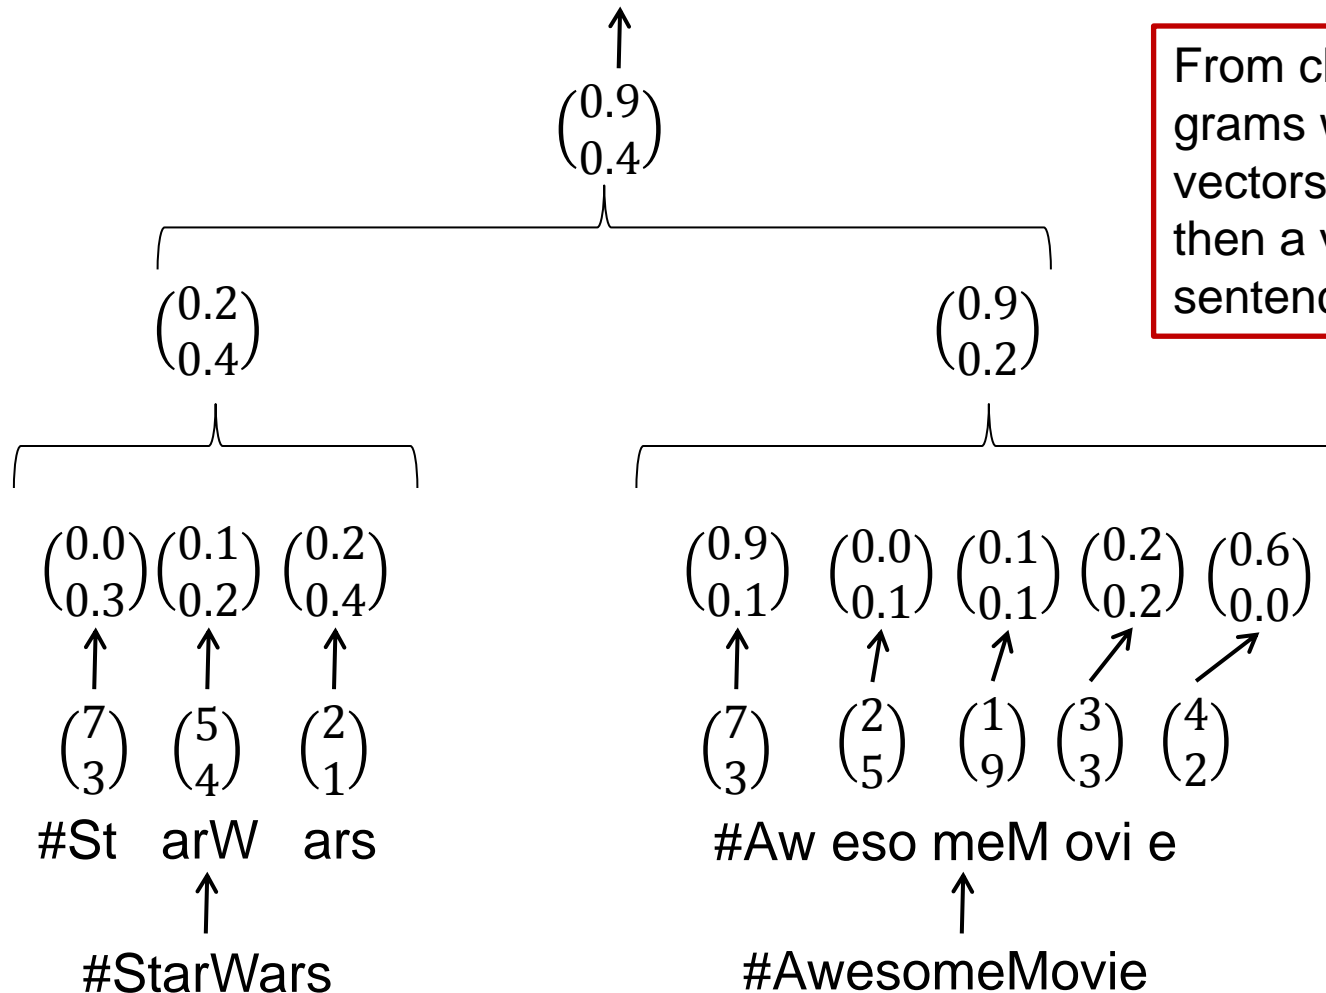
- We can create convolutional layers working on unigrams, bigrams, trigrams etc. and combine their output



# Or work on a different granularity



# Stacking Convolutional Layers

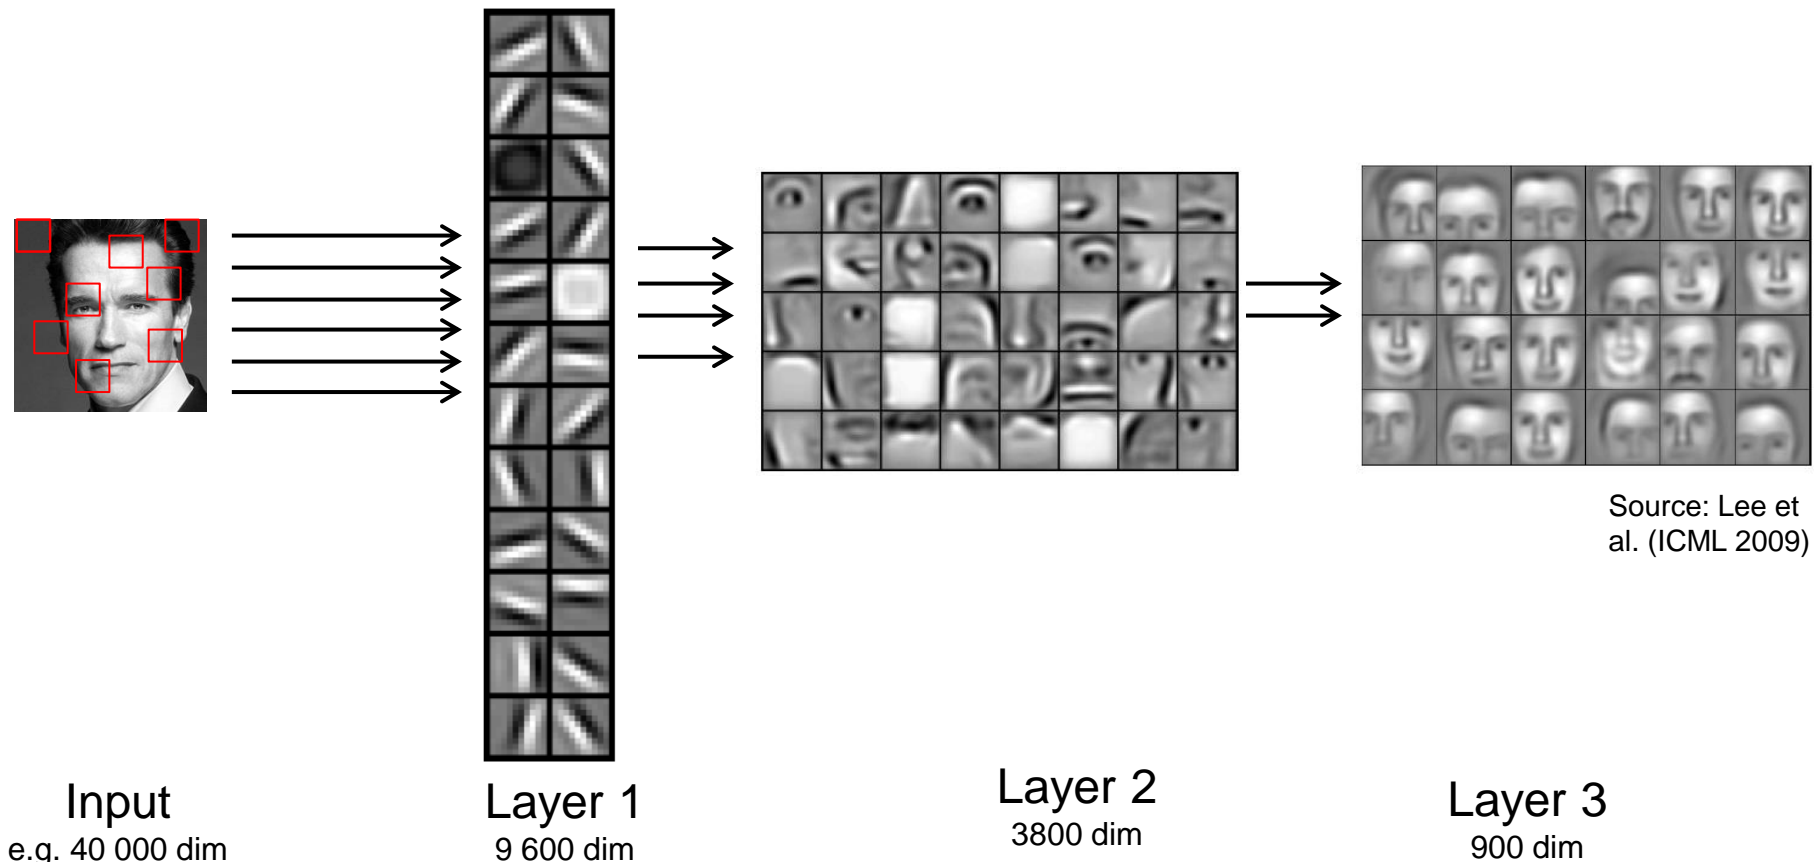


From character n-grams we could derive vectors for words and then a vector for the sentence



# Stacked Convolutional Layers

- Computer vision uses stacked convolutional layers to derive from simple representations high level representations

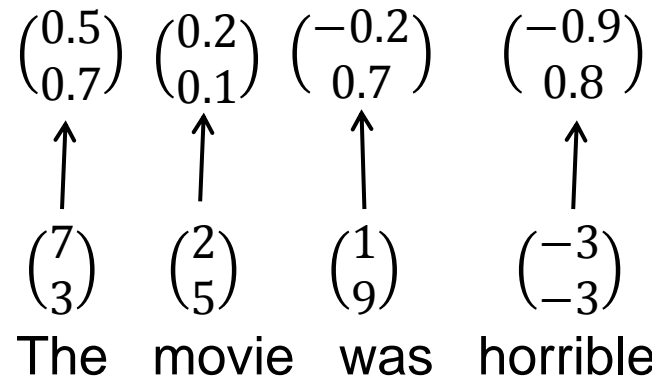


Source: Lee et al. (ICML 2009)

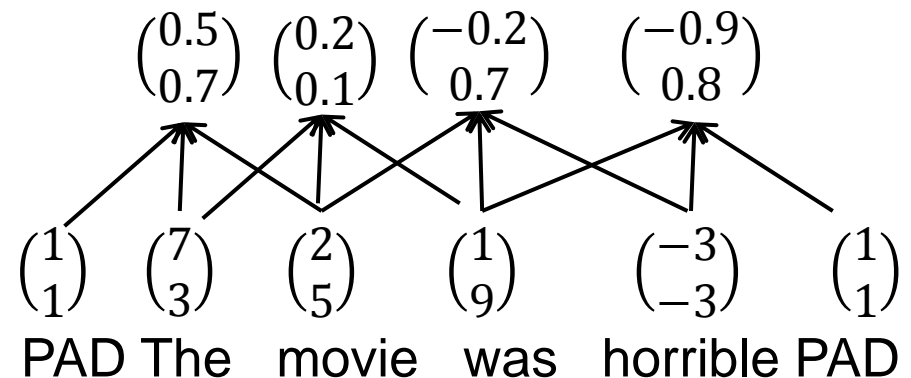
# Terminology: Filter Length

- The filter length is the extension of each filter
- Mainly inspired by Computer Vision where we work on spatial close pixels
- In NLP we are more flexible:
  - Use a context window of size  $n$
  - Use dependency links / syntax tree

Filter Length: 1

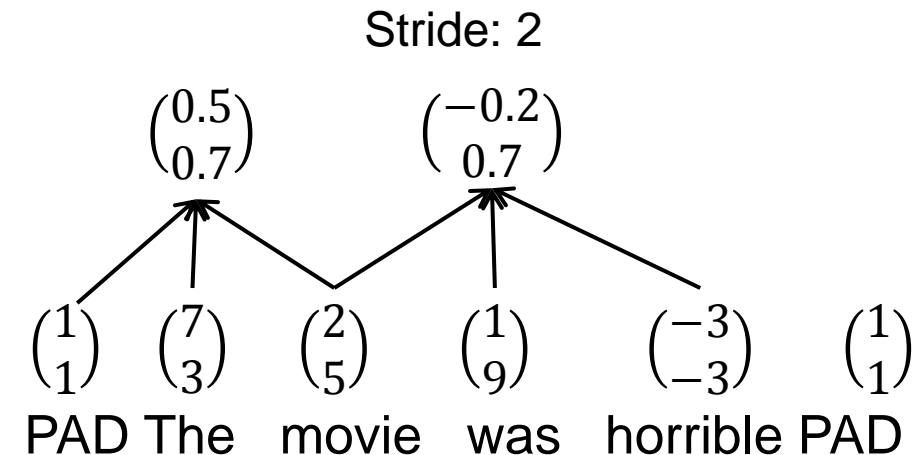
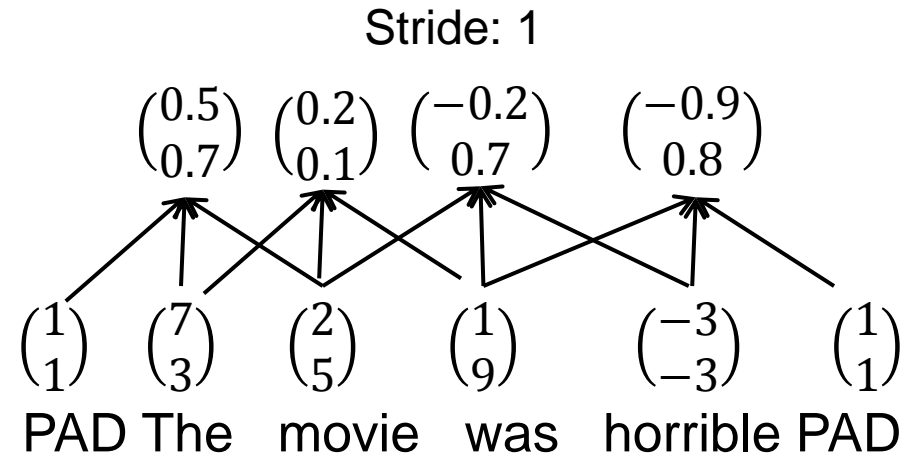


Filter Length: 3



# Terminology: Stride

- The *stride* specifies the steps size we move across a sentence
- In NLP: Typically stride=1
- In computer vision: Other values can be used



# How to choose the embeddings?

- When we use words, how should we initialize the embeddings?

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-rand	76.1	45.0	82.7	89.6	91.2	79.8	83.4
CNN-static	81.0	45.5	86.8	93.0	92.8	84.7	<b>89.6</b>
CNN-non-static	<b>81.5</b>	48.0	87.2	93.4	93.6	84.3	89.5
CNN-multichannel	81.1	47.4	<b>88.1</b>	93.2	92.2	<b>85.0</b>	89.4

- Rand: Random initialization
- Static: word2vec, no updates during training
- Non-static: word2vec with updates
- Multichannel: next slide

Source: Kim, 2014. Performance on Sentence Classification Tasks

# Multi-Channel Idea

- We start with 2 copies for the word vectors, both initialized with word2vec/GloVe etc.
- Only one version of them is updated, the other is static
- We apply the same filters to both channels before we apply the pooling
- The one channel can learn task specific embeddings
  - E.g. for sentiment, *good* and *bad* should be far away in vector space
- So far mixed results
- Different Idea: Use differently pre-trained word embeddings
  - E.g. based on local context, on dependency trees, on relations from knowledge bases etc.
- Reference: Kim, 2014. *Performance on Sentence Classification Tasks*

# Hints on the Implementation

- Numpy and Theano cannot work with variable sized rows
- How to model a dataset like this?  
[ [This is my first sentence .]  
[This is my second , longer sentence .]  
[Super short ]]
- 2 Approaches:
  - Ignore minibatches, just input 1 sentence at a time for training / testing
    - Bad for performance
  - Pad the sentences with 0 to make them of the same length
    - Be careful with the padding, that the max-pooling does not choose your padded values
    - Be careful with the runtime, that a single super long sentence does not create too much padding for all other sentences.
    - Great to run on GPU (convolutions can be computed easily in parallel)

# Hints on the Implementation II

- Most implementation for convolutional layers are targeted for computer vision
- They introduce a lot more hyper parameters.
- Keras.Convolution1D:
  - nb\_filter: dimensionality of the output
  - filter\_length: The extension (spatial or temporal) of each filter.
  - border\_mode: How are vales at the border handled. 'valid' or 'full'.
  - subsample\_length: The stride value for the filter
- Keras.MaxPooling1D:
  - pool\_length: factor by which to downscale. 2 will halve the input.
  - stride: Stride value.
  - ignore\_border: boolean
- **Note:** Existent Convolution1D and MaxPooling1D not suitable for max-over-time implementation